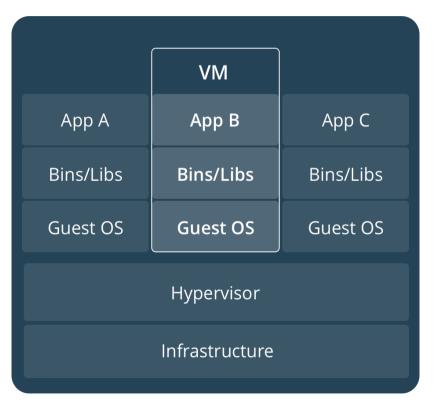


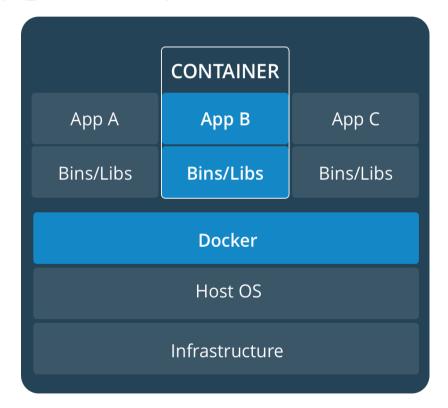
Radu Munteanu – 1 November 2018

Contents

- Why Kubernetes
- Why Extend
- Kubernetes Architecture
- What to Extend
- How to Extend
- Demo

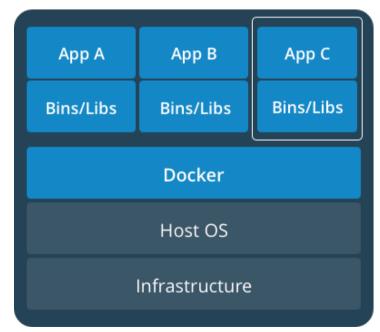
Virtual Machines vs Containers (docker)

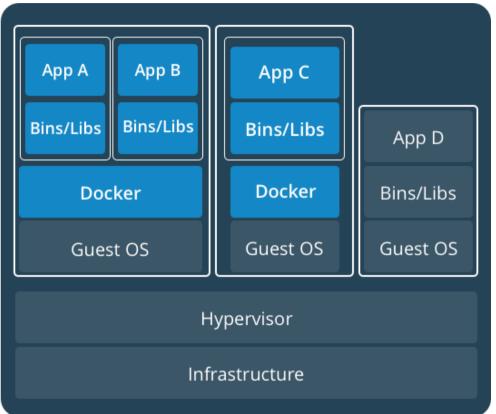




https://www.docker.com/what-container

Containers on VMs

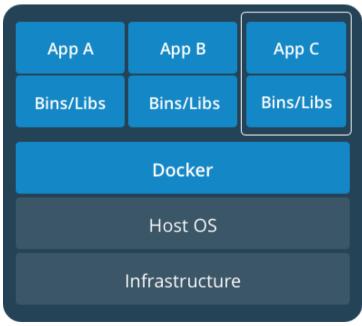


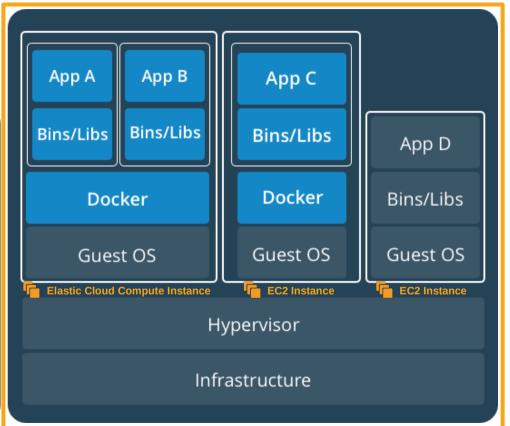


https://www.docker.com/what-container



Containers on VMs





https://www.docker.com/what-container

- W Kubernetes (K8s)
 - Container communication over multiple hosts
 - Higher level of abstraction: resources
 (E.g. nodes, pods, replica sets, deployments, services)
 - => Easier definition of virtual resources (VMs, apps, proxies)
 - => Automatic management of failures (replica sets, deployments)
 - Central management of containers (Dashboard or kubectl)
 - One declarative, domain-specific language for managing all resources within the Kubernetes cluster (common format: YAML)

Why Extend

- Easier to manage the whole infrastructure using the same language and tools
 - Managing AWS resources
 - Now: AWS Console → Cloud Formation DSL, aws-cli, aws-api
 - Future: Cloud Formation DSL, aws-cli, aws-api → K8s DSL, kubectl
 - Managing installation and deployments
 - aws-cli, scripts → K8s DSL, kubectl
- Standardize communication and configuration across all apps using the features and extensibility of K8s
 - Service Discovery: Zookeeper → K8s Service Catalog
 - Configuration: Zookeeper → Custom Resources, Custom K8s API

- Resources (Objects *)
 - A Kubernetes resource is a declarative API with a well defined **Schema** structure and **endpoints**. Because the structure of the Schema and Endpoints are both well understood, many Kubernetes tools support all APIs written as Kubernetes resources
 - Schema: Group, Version, Kind

Every Kubernetes resource has a Group, Version and Kind that uniquely identifies it

- The resource **Kind** is the name of the API such as Deployment or Service
- The resource Version defines the stability of the API and backward compatibility guarantees such as v1beta1 or v1
- The resource **Group** is similar to package in a language. It disambiguates different APIs that may happen to have identically named Kinds. Groups often contain a domain name, such as k8s.io

apiVersion: apps/v1 apiVersion: radu/v1alpha1

kind: Deployment kind: MyResource

* https://kubernetes.io/docs/concepts/overview/working-with-objects/kubernetes-objects http://book.kubebuilder.io/basics/what_is_a_resource.html

Resources

Schema: Spec, Status, Metadata

Most Kubernetes resources Schema contain 3 components: Spec, Status and Metadata

- **Spec**: defines the desired state of the cluster as specified by the user
- **Status**: publishes the state of the cluster as observed by the controller
- Metadata: contains information common to most resources about the object including as the object name, annotations, labels and more

Resources

```
apiVersion: extensions/v1beta1
kind: Deployment
metadata:
  name: nginx
  namespace: default
spec:
  replicas: 1
  template:
    spec:
      containers:
      - image: nginx
        name: nginx
status:
  replicas: 1
  unavailableReplicas: 1
  updatedReplicas: 1
```

• Resources

Endpoints

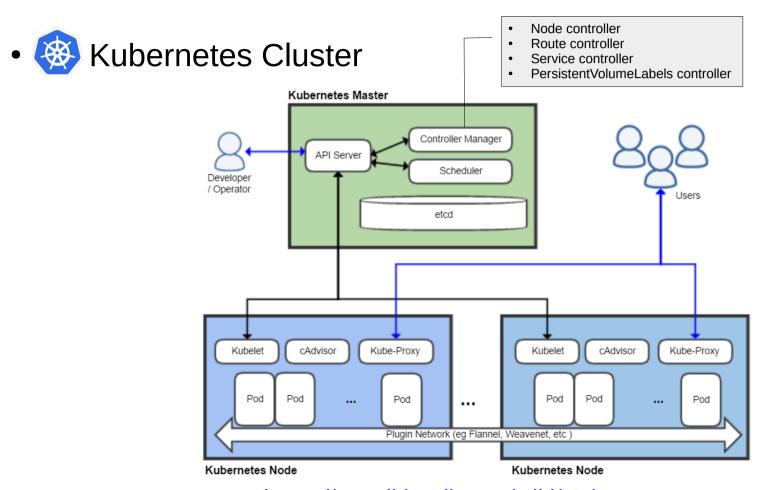
Kubernetes resources have well defined endpoints as described below

- Create, Update, Patch, Delete
 - The create, update, patch and delete endpoints may be used to modify objects. The update endpoint replaces the object with what is provided, whereas the patch endpoint selectively updates fields
- Get, List, Watch

The get, list and watch endpoints may be used to get a specific resource by name, list all resources matching a labels, or continually watch for updates

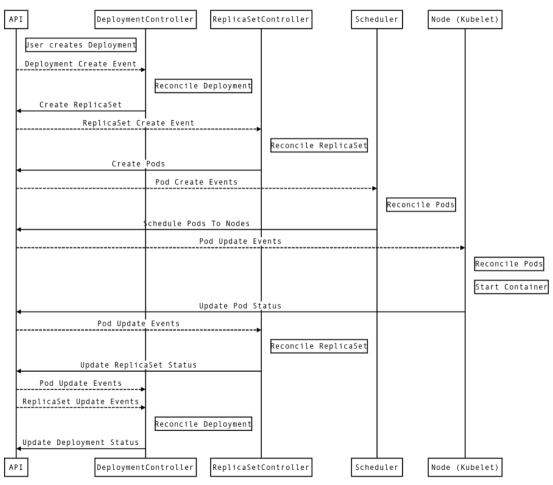
• **®** Namespaces

- Kubernetes supports multiple virtual clusters backed by the same physical cluster. These virtual clusters are called namespaces
 - Namespaces are intended for use in environments with many users spread across multiple teams, or projects. For clusters with a few to tens of users, you should not need to create or think about namespaces at all. Start using namespaces when you need the features they provide
 - Namespaces provide a scope for names. Names of resources need to be unique within a namespace, but not across namespaces
 - Namespaces are a way to divide cluster resources between multiple users (via resource quota)



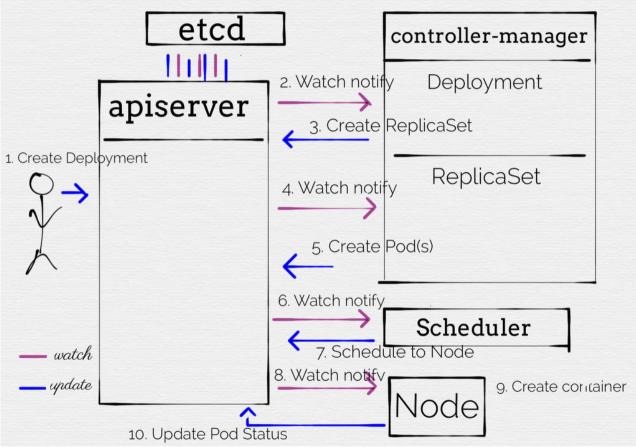
https://en.wikipedia.org/wiki/Kubernetes





http://book.kubebuilder.io/basics/what_is_a_controller.html





What to Extend

- Kubectl Plugins
- Service Catalog
- Custom Resources
- Custom Schedulers
- Custom Controllers
- Custom APIs
- Network Plugins
- Storage Plugins
- ...

How to Extend – Kubectl Plugins

Calling a Plugin

```
kubectl plugin <plugin name> [flags] [options]
```

Plugin Structure

```
~/.kube/plugins/
Largaryen
Plugin.yaml
dracarys
```

Plugin YAML

```
name: "targaryen"
shortDesc: "Dragonized plugin"
longDesc: ""
example: ""
command: "./dracarys"
flags:
   - name: "heat"
    shorthand: "h"
    desc: "Fire heat"
    defValue: "extreme"
tree:
```

Environment Variables

```
KUBECTL PLUGINS
```

```
# REQUIRED: the plugin command name, to be invoked under 'kubectl'
# REQUIRED: the command short description, for help
# the command long description, for help
# command example(s), for help
# REQUIRED: the command, binary, or script to invoke
# flags supported by the plugin
# REQUIRED for each flag: flag name
# short version of the flag name
# REQUIRED for each flag: flag description
# default value of the flag
# allows the declaration of subcommands
# subcommands support the same set of attributes
```

https://v1-11.docs.kubernetes.io/docs/tasks/extend-kubectl/kubectl-plugins

How to Extend – Kubectl Plugins

- In V.1.12
- Calling a Plugin

```
kubectl <plugin_name> [flags] [options]
```

- Plugin Structure
 - No more structure, the plugin is any app found in PATH starting with kubectl-

```
kubectl-<plugin name>
```

- Subcommands are defined by adding a dash in the name

```
kubectl-<plugin name>-<subcommand1>-<subcommand2>
```

- E.g. kubectl foo bar baz → kubectl-foo-bar-baz
- Define a plugin with the name hello-world?

```
https://kubernetes.io/docs/tasks/extend-kubectl/kubectl-plugins
```

Custom Resource Definition (CRD)

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  # name must match the spec fields below, and be in the form: <plural>.<group>
  name: crontabs.stable.example.com
spec:
  # group name to use for REST API: /apis/<group>/<version>
  group: stable.example.com
  # version name to use for REST API: /apis/<group>/<version>
  version: v1
  # either Namespaced or Cluster
  scope: Namespaced
  names:
    # plural name to be used in the URL: /apis/<group>/<version>/<plural>
   plural: crontabs
    # singular name to be used as an alias on the CLI and for display
    singular: crontab
    # kind is normally the CamelCased singular type. Your resource manifests use this.
    kind: CronTab
    # shortNames allow shorter string to match your resource on the CLI
    shortNames:
    - ct
```

https://kubernetes.io/docs/tasks/access-kubernetes-api/extend-api-custom-resource-definitions

CRD

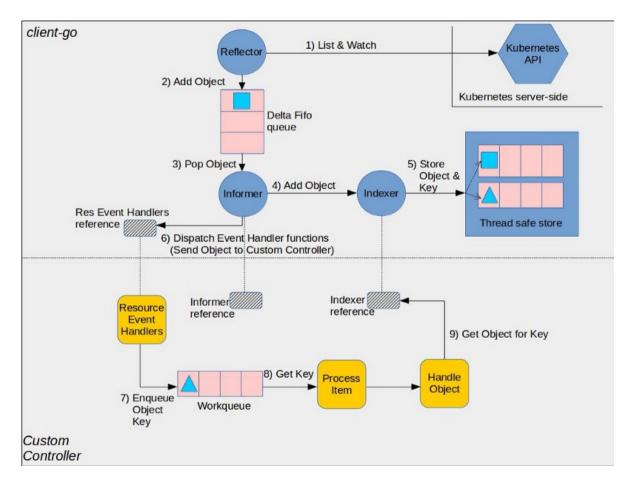
```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  # name must match the spec fields below, and be in the form: <plural>.<group>
  name: crontabs.stable.example.com
spec:
  # group name to use for REST API: /apis/<group>/<version>
  group: stable.example.com
  # list of versions supported by this CustomResourceDefinition
  versions:
    - name: v1
      # Each version can be enabled/disabled by Served flag.
      served: true
      # One and only one version must be marked as the storage version.
      storage: true
  # either Namespaced or Cluster
  scope: Namespaced
  names:
    # plural name to be used in the URL: /apis/<group>/<version>/<plural>
    plural: crontabs
    # singular name to be used as an alias on the CLI and for display
    singular: crontab
    # kind is normally the CamelCased singular type. Your resource manifests use this.
    kind: CronTab
    # shortNames allow shorter string to match your resource on the CLI
    shortNames:
```

https://kubernetes.io/docs/tasks/access-kubernetes-api/extend-api-custom-resource-definitions

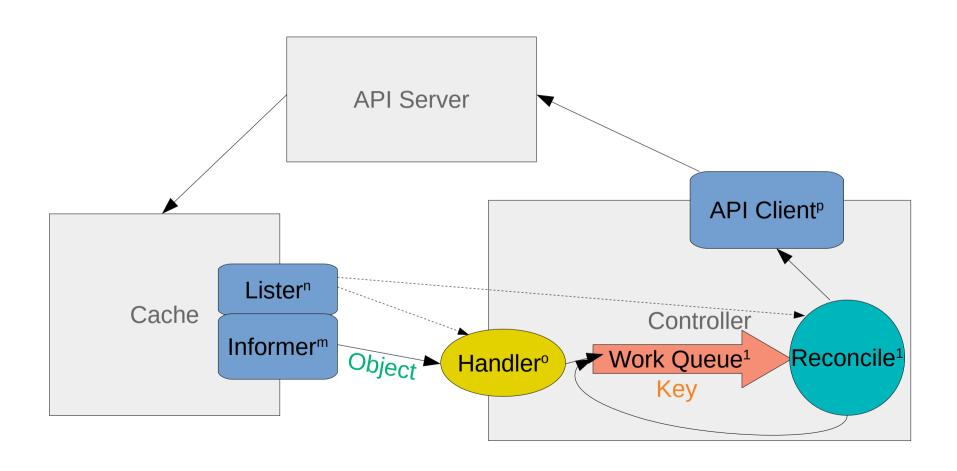
- Why?
 - Automate your infrastructure through Kubernetes DSL
 - Interact with Operators* for your Cloud Native Apps

* "Some controllers are referred to as Operators. Operators are a specific type of controller that manage running a specific application such as Redis or Cassandra." (http://book.kubebuilder.io/basics/what is a controller.html)

How to use in Controllers?



https://medium.com/@cloudark/kubernetes-custom-controllers-b6c7d0668fdf



- Leader Election
 - "k8s.io/client-go/tools/leaderelection"
 - https://github.com/kubernetes/client-go/tree/master/tools/leaderelection
 - Resource Locks:
 - Endpoints
 - ConfigMaps

- There's no mature library or framework
- What can you use right now:
 - Modifying existing controllers (like sample-controller)
 - KubeBuilder
 - MetaController
 - Operator SDK
 - k8s.io Code-Generator & co (client-go, apimachinery, etc.)
 - sigs.k8s.io controller-runtime, controller-tools

Demo

- sample-controller
 - https://github.com/kubernetes/sample-controller
- k8s-sample-plugin
 - https://gitlab.com/radu-munteanu/k8s-sample-plugin

Q & A